Chapter 3 Transport Layer

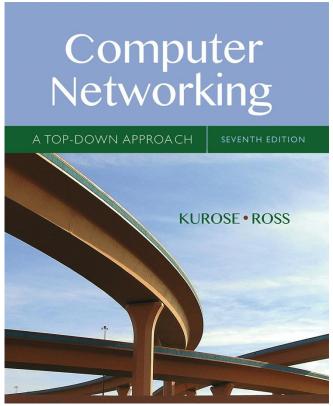
A note on the use of these Powerpoint slides:

We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

Thanks and enjoy! JFK/KWR

© All material copyright 1996-2016 J.F Kurose and K.W. Ross, All Rights Reserved



Computer Networking: A Top Down Approach

7th edition
Jim Kurose, Keith Ross
Pearson/Addison Wesley
April 2016

Transport Layer

Chapter 3: Transport Layer

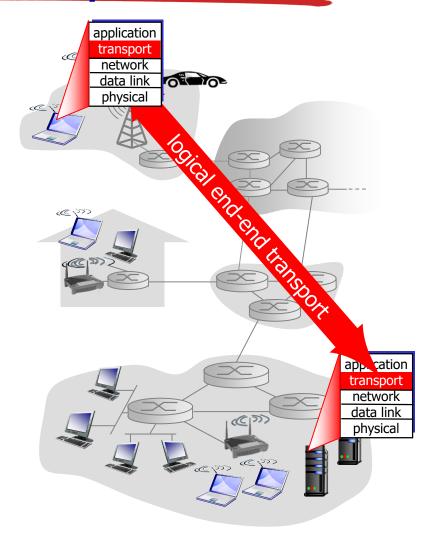
our goals:

- understand principles behind transport layer services:
 - multiplexing, demultiplexing
 - reliable data transfer
 - flow control
 - congestion control

- learn about Internet transport layer protocols:
 - UDP: connectionless transport
 - TCP: connection-oriented reliable transport
 - TCP congestion control

Transport services and protocols

- provide logical communication between app processes running on different hosts
- transport protocols run in end systems
 - send side: breaks app messages into segments, passes to network layer
 - rcv side: reassembles segments into messages, passes to app layer
- more than one transport protocol available to apps
 - Internet: TCP and UDP



Transport vs. network layer


- network layer: logical communication between hosts (IP address to IP address)
- transport layer: logical communication between processes (port # to port #)
 - relies on network layer services

household analogy:

- 12 kids in Ann's house sending letters to 12 kids in Bill's house:
- hosts = houses
- apps = kids
- messages = letters in envelopes
- transport protocol = Ann and Bill who demux to inhouse siblings
- network-layer protocol = postal service


Internet transport-layer protocols

- reliable, in-order delivery (TCP)
 - congestion control
 - flow control
 - connection setup
- unreliable, unordered delivery: UDP
 - no-frills extension of "best-effort" IP
- services not guaranteed:
 - no delay guarantees
 - no bandwidth guarantees




Protocol Multiplexing

Multiplexing/demultiplexing



How demultiplexing works

- host receives IP datagrams
 - each datagram has source IP address, destination IP address
 - each datagram carries one transport-layer segment
 - each segment has source and destination port #
- host uses IP addresses & port numbers to direct segment to appropriate socket



TCP/UDP segment format

Connectionless demultiplexing

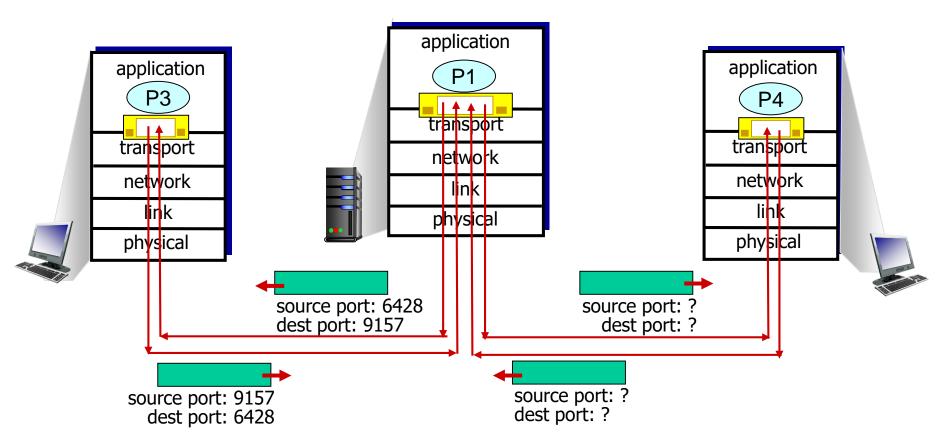
- recall: when sending datagrams into UDP socket, we must specify
 - destination IP address
 - destination port #

- when dest. host receives UDP segment:
 - checks destination port # in segment



 directs UDP segment to socket with that port # IP datagrams with same dest. port #, but different source IP addresses (and/or source port #) will be directed to same socket at destination host

Connectionless demux: example




Connection-oriented demux

- TCP socket identified by 4-tuple (Why?):
 - source IP address
 - source port number
 - dest IP address
 - dest port number
- demux: receiver uses all four values to direct segment to appropriate socket

- server host may support many simultaneous TCP sockets:
 - each socket identified by its own 4-tuple
- web servers have different sockets for each connecting client
 - non-persistent HTTP will have different socket for each request

Connection-oriented demux: example



three segments, all destined to IP address: B, dest port: 80 are demultiplexed to *different* sockets

UDP

UDP: User Datagram Protocol [RFC 768]

- "no frills", "bare bones" Internet transport protocol
- "best effort" service, UDP segments may be:
 - lost
 - delivered out-of-order to app
- connectionless:
 - no handshaking between UDP sender, receiver
 - each UDP segment is handled independently

- UDP use:
 - streaming multimedia apps (loss tolerant, rate sensitive)
 - DNS
 - SNMP
- reliable transfer over UDP (e.g., QUIC):
 - add reliability at application layer
 - application-specific error recovery!

UDP: segment header

32 bits dest port # source port # checksum length application data (payload)

UDP segment format

length, in bytes of UDP segment, including header

why is there a UDP?

- no connection establishment (which can add delay)
- simple: no state at sender or receiver
- small header size
- no congestion control: UDP can blast away as fast as desired

UDP checksum

Goal: detect "errors" (e.g., flipped bits) in transmitted segment

sender:

- treat segment contents, including header fields, as sequence of 16-bit integers
- checksum: addition (one's complement sum) of segment contents
- sender puts checksum value into UDP checksum field


receiver:

- compute checksum of received segment
- check if computed checksum equals checksum field value:
 - NO error detected
 - YES no error detected.
 But maybe errors
 nonetheless? More later

. . . .

Internet checksum: example

example: add two 16-bit integers



Note: when adding numbers, an overflow bit from the most significant bit needs to be added to the result

^{*} Check out the online interactive exercises for more examples: http://gaia.cs.umass.edu/kurose_ross/interactive/

Internet checksum: generation

- I. initial cs = 0;
- 2. sum = sum all 16-bit words in segment including optional padding and initial cs;
- 3. cs = one's complement of sum
- 4. store cs in the checksum field before sending
- Note: cs is NEVER 0; if cs = 0 then store FFFF

Internet checksum: validation

- sum all 16-bit words including optional padding and checksum (which is NEVER 0).
- let sum' = one's complement of sum
- Note:


```
sum + ~sum = FFFF
FFFF + FFFF = FFFF (add overflow bit)
0000 = ~FFFF.
```

• if sum' = 0000 then checksum accepted otherwise error

Reliable Data Transfer

Principles of reliable data transfer

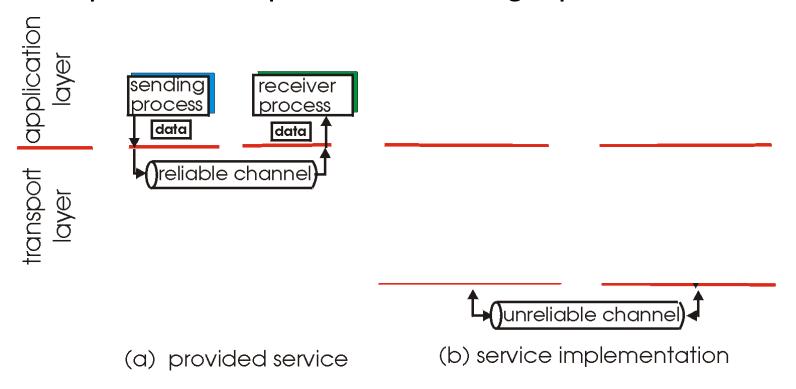
- important in application, transport, link layers
 - top-I0 list of important networking topics!



- (a) provided service
- characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

Principles of reliable data transfer


- important in application, transport, link layers
 - top-10 list of important networking topics!



 characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)


Principles of reliable data transfer

- important in application, transport, link layers
 - top-10 list of important networking topics!



 characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

Reliable data transfer: events



Reliable data transfer: getting started

- consider only unidirectional data transfer
 - but control info will flow in both directions!
- use finite state machines (FSM) to specify sender, receiver

state: when in this "state" next state uniquely determined by next event and condition

event causing state transition [condition]

actions taken on state transition

event [condition]

actions

event causing state transition [condition]


actions taken on state transition

event [condition]

rdt: channel with bit errors

- underlying channel may flip bits in packet
 - checksum to detect bit errors
- the question: how to recover from errors:
 - acknowledgements (ACKs): receiver explicitly tells sender that packet received OK
 - negative acknowledgements (NAKs): receiver explicitly tells sender that packet had errors
 - sender retransmits packet on receipt of NAKs
- new mechanisms:
 - error detection
 - feedback: control messages (ACKs, NAKs) from receiver to sender

rdt (I): specification




sender

receiver

recv(data,cs) [cs error]





recv(data,cs) [cs ok]

deliver(data); send(ACK)

rdt (I) has a fatal flaw!


what happens if ACK/NAK corrupted or lost?

- sender doesn't know what happened at receiver!
- can't just retransmit: possible duplicate


handling duplicates:

- sender retransmits current packet if ACK/NAK corrupted
- sender adds sequence number to each packet
- receiver discards (doesn't deliver up) duplicates

rdt (2): ACK with sequence # (sender)



rdt (2): ACK with sequence # (receiver)




rdt (3): error, duplicate, and loss


sender waits "reasonable" amount of time for ACK:

- retransmits if timeout
- retransmission will be duplicated if packet is delayed instead of lost
- the received ACK# must match, what is expected (i.e., ACK# means ``receiver seen packet #"); retransmit if not
- known as "Alternating Bit Protocol" (ABP).


rdt (3): ACK, sequence #, timer (sender)




rdt (3) in action



rdt (3) in action





(d) premature timeout/ delayed ACK

Performance of rdt (3)

- rdt (3) is correct, but performance stinks!
- e.g.: I Gbps link, 30 ms RTT delay, 8000 bit packet:


$$D_{trans} = \frac{L}{R} = \frac{8000 \text{ bits}}{10^9 \text{ bits/sec}} = 8 \text{ microsecs}$$

U sender: utilization – fraction of time sender busy sending

$$U_{\text{sender}} = \frac{L/R}{RTT + L/R} = \frac{.008}{30.008} = 0.00027$$


- if RTT=30 msec, IKB packet every 30 msec: 33kB/sec throughtput over I Gbps link
- network protocol limits use of physical resources!

rdt (3): stop-and-wait operation



Sliding Window Protocols

Pipelining: increased utilization




Pipeline capacity

- It is defined as bandwidth-delay product.
- For example, bandwidth = I Mbps (I million bits/sec), and RTT delay = I second.
- Then, the pipeline capacity is I Mbps x I second = I million bits.
- The higher the bandwidth-delay product; the larger the pipeline capacity.
- However, network "capacity" varies over time, due to congestion and packet loss.

Pipelined protocols

pipelining: sender allows multiple, "in-flight", yet-to-be-acknowledged packets

- range of sequence numbers must be increased
- buffering at sender and/or receiver



(b) a pipelined protocol in operation

two generic forms of pipelined protocols: go-Back-N, selective repeat

Sliding window protocols

Go-back-N:

- sender can have up to N unack'ed packets in pipeline
- receiver only sends cumulative ack, what it has seen so far
- sender has timer for oldest unack'ed packet
 - when timer expires, retransmit all unacked packets


Selective Repeat:

- sender can have up to N unack'ed packets in pipeline
- receiver sends individual ack for each packet

- sender maintains a timer for each unack'ed packet
 - when a timer expires, retransmit only that unack'ed packet

Go-Back-N: sender

- assign a k-bit seq # in packet header (2^k > N)
- "window" of up to N, consecutive unack'ed packets allowed




- ACK(n): all packets up to and including seq #n received "cumulative ACK"
 - sender may receive duplicate ACKs (see receiver)
- a timer for the oldest (earliest) unack'ed packet
- timeout(n): retransmit all unack'ed packets in window

Go-Back-N: receiver

- Use the same sequence #s as the sender
- But, it has a window of size I
- The window specifies the next expected packet sequence #.
- Cumulative ACKs: ACK(n) if all packets up to and including sequence #n have been received.
- Out of order packets are discarded!

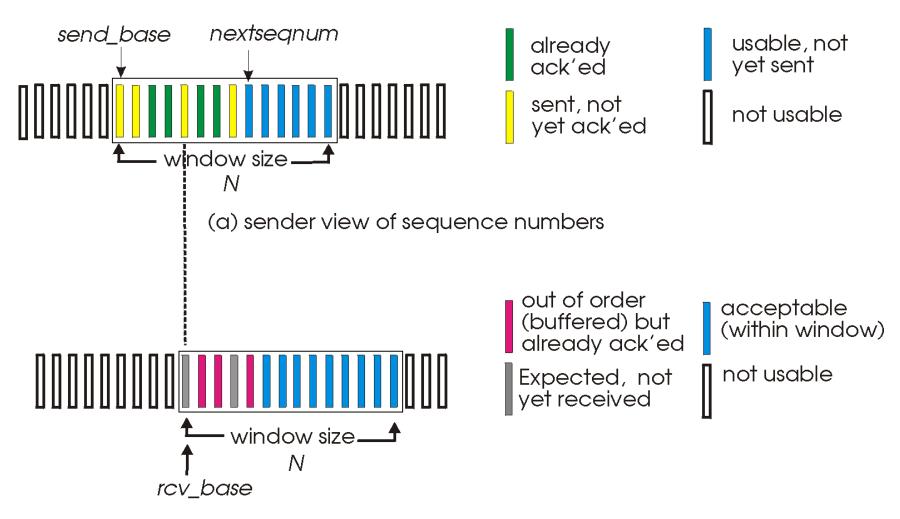
GBN in action



Selective repeat

- receiver individually acknowledges all correctly received packets within window
 - buffer packets, as needed, for eventual in-order delivery to upper layer
- sender only resends packets for which ACK not received (i.e., timeout)
 - sender set timer for each unACK'ed packet
- sender window
 - windows size N <= half of max. sequence #
 - limits seq #s of sent, unACK'ed packets
 - advance window when consecutive packets have been ACK'ed

Selective repeat: sender, receiver windows



(b) receiver view of sequence numbers

Selective repeat

sender

data from above:

if next available seq # in window, send pkt

timeout(n):

resend pkt n, restart timer

ACK(n) in window:

- mark pkt n as received
- If m smallest unACKed pkt #, advance window base to m

receiver -

pkt n in window:

- send ACK(n)
- out-of-order: buffer
- in-order: deliver all inorder consecutive pkts, advance window to next not-yet-received pkt

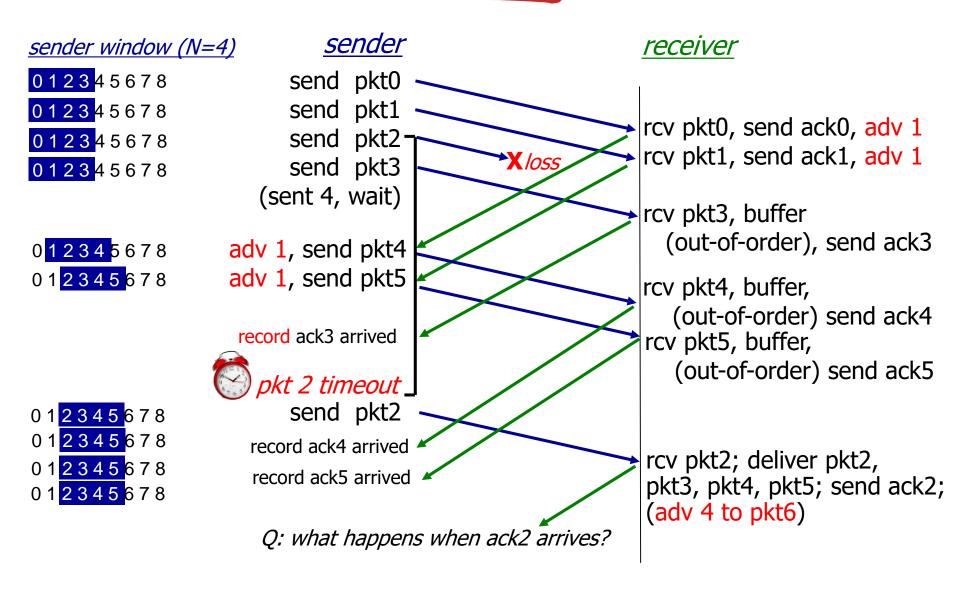
pkt n was seen:

ACK(n)

otherwise:

ignore

Selective repeat in action

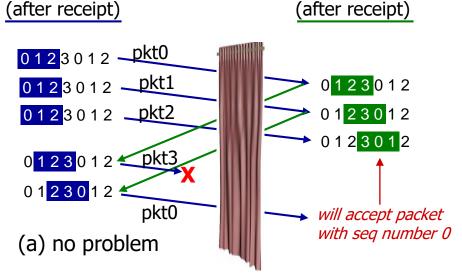


Selective repeat: dilemma

example:

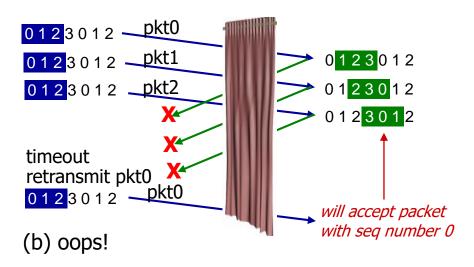
- seq #' s: 0, 1, 2, 3
- window size=3
- receiver sees no difference in two scenarios!
- duplicate data accepted as new in (b)
- Q: what relationship between seq # size and window size N to avoid problem in (b)?

N <= half of seq #



sender window

receiver can't see sender side.
receiver behavior identical in both cases!
something's (very) wrong!



receiver window

TCP

TCP: Overview

- point-to-point:
 - one sender, one receiver
- reliable, in-order byte steam:
 - no "message boundaries"
- pipelined:
 - TCP congestion and flow control set window size automatically and dynamically

full duplex data:

- bi-directional data flow in same connection
- MSS: maximum segment size (< MTU)
- connection-oriented:
 - handshaking (3-way)
 agree on sender's, and
 receiver's initial
 sequence # and window
 sizes

TCP segment structure

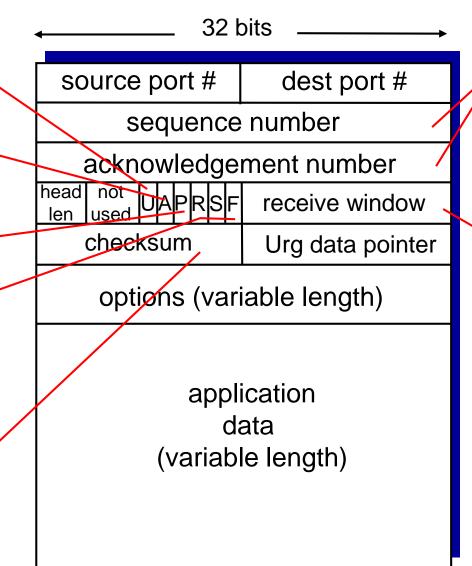
URG: urgent data (generally not used)

ACK: ACK # valid

PSH: push data now (generally not used)

RST, SYN, FIN: connection estab (setup, teardown commands)

> Internet checksum' (as in UDP)



by bytes of data (not segments!)

bytes
rcvr willing
to accept


TCP seq. numbers, ACKs

sequence numbers:


 byte stream "number" of first byte in segment's data

acknowledgements:

- seq # of next byte expected from other side
- cumulative ACK
- Q: how receiver handles out-of-order segments
 - A: TCP spec doesn't say,
 - up to implementor



TCP seq. numbers, ACKs



simple telnet scenario

TCP round trip time, timeout

- Q: how to set TCP timeout value?
- longer than RTT
 - but RTT varies
- too short: premature timeout, unnecessary retransmissions
- too long: slow reaction to segment loss


Q: how to estimate RTT?

- SampleRTT: measured time from segment transmission until ACK receipt
 - ignore retransmissions
- SampleRTT will vary, use exponential moving average
 - estimatedRTT: average several recent measurements, not just current SampleRTT

TCP round trip time, timeout

EstimatedRTT = $(1-\alpha)$ *EstimatedRTT + α *SampleRTT

- exponential weighted moving average
- influence of past sample decreases exponentially fast
- typical value: $\alpha = 0.125$ (less sensitive to SampleRTT)



TCP round trip time, timeout

- timeout interval: EstimatedRTT plus "safety margin"
- estimate SampleRTT deviation from EstimatedRTT:

```
DevRTT = (1-\beta)*DevRTT + \beta*|SampleRTT-EstimatedRTT| (typically, \beta = 0.25)
```

TCP reliable data transfer

- TCP creates reliable data service on top of IP's unreliable service
 - pipelined (sliding window) segments
 - cumulative ACKs
 - single retransmission timer
- retransmissions triggered by:
 - timeout events
 - 3 duplicate ACKs

TCP sender events:


data (from app):

- create segment with seq #
- seq # is byte-stream number of first data byte in segment
- start timer if not already running
 - think of timer for oldest unACK'ed segment
 - expiration interval:TimeOutInterval


Timeout (internal):

- retransmit segment that caused timeout
- restart timer
 ack (from receiver):
- if ACK acknowledges previously unacked segments
 - update what is known to be ACK'ed
 - start timer if there are still unACK'ed segments

TCP: retransmission scenarios



TCP: retransmission scenarios



cumulative ACK

TCP fast retransmit


- time-out period often relatively long:
 - long delay before resending lost packet
- detect lost segments via duplicate ACKs.
 - sender often sends many segments backto-back
 - if segment is lost, there will likely be many duplicate ACKs.

TCP fast retransmit

if sender receives 3
ACKs for same data,
resend unACK'ed
segment with smallest
seq #

likely that unACK'ed segment is lost, so don't wait for timeout

TCP fast retransmit



fast retransmit after sender receipt of triple duplicate ACKs (total 4)

TCP Flow Control

TCP flow control

application may remove data from TCP socket buffers

... slower than TCP receiver is delivering (sender is sending)


application process application OS TCP socket receiver buffers TCP code ĬΡ code from sender

receiver protocol stack

receiver controls sender, so sender won't overflow receiver's buffer by transmitting too much, too fast

TCP flow control

- receiver "advertises" free buffer space by including rwnd value in TCP header of receiver-to-sender segments
 - RcvBuffer size set via socket options (typical default is 4096 bytes)
 - many operating systems autoadjust RcvBuffer
- sender limits amount of unACK'ed data to receiver's rwnd value
- sender stops sending if rwnd == 0.




receiver-side buffering

Connection Management

before exchanging data, sender/receiver "handshake":

- agree to establish connection (each knowing the other willing to establish connection)
- agree on connection parameters (initial seq. #s, rcvBuffer size)



```
application

connection state: ESTAB
connection Variables:
  seq # client-to-server
        server-to-client
   rcvBuffer size
        at server,client

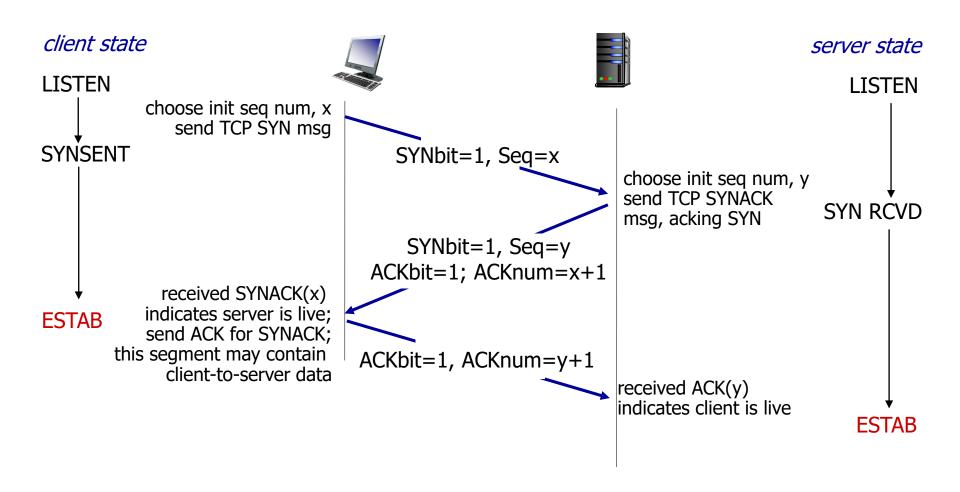
network
```

```
Socket clientSocket =
  newSocket("hostname","port
  number");
```

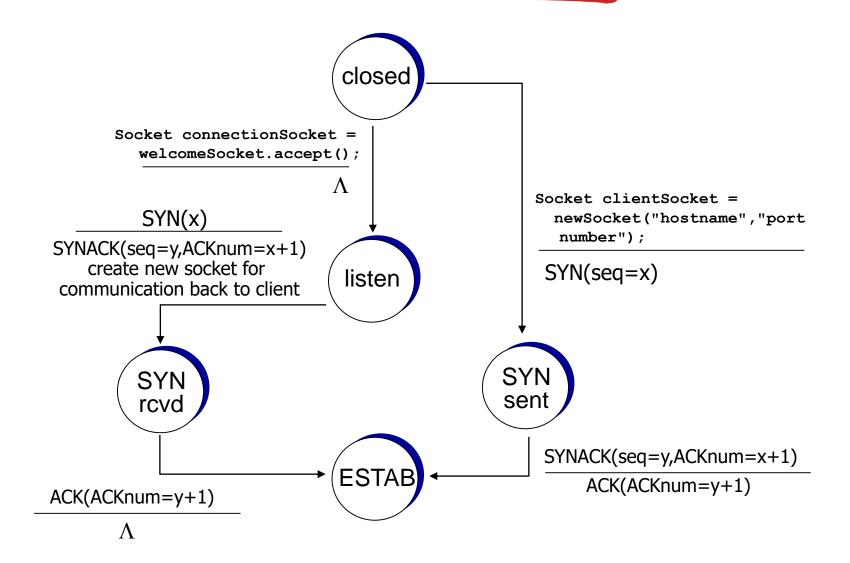
```
Socket connectionSocket =
  welcomeSocket.accept();
```

TCP Connection Setup

TCP 3-way handshake



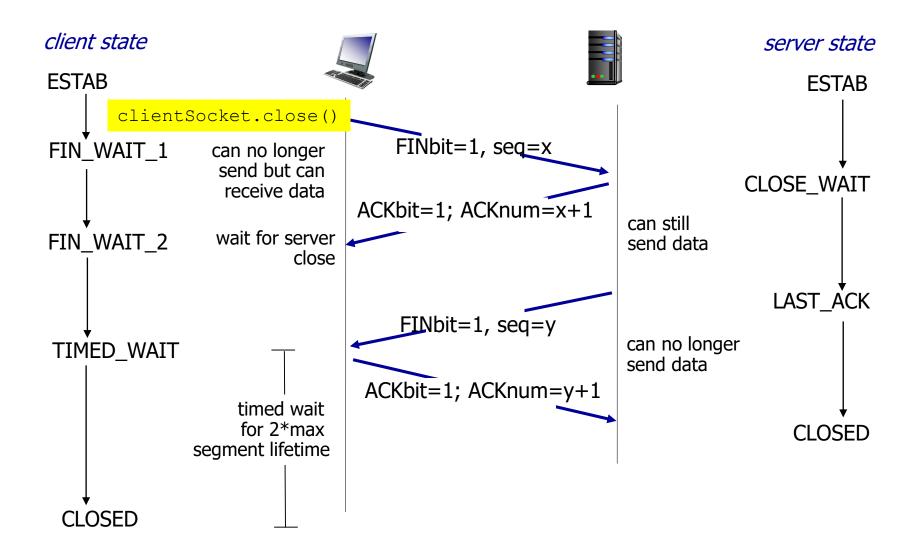
TCP 3-way handshake: FSM



TCP: closing a connection

- client, server each close their side of connection
 - send TCP segment with FIN bit = I
- respond to received FIN with ACK
 - on receiving FIN, ACK can be combined with own FIN
- simultaneous FIN exchanges can be handled

TCP: closing a connection



TCP Congestion Control


Principles of congestion control

- informally: "too many sources sending too much data too fast for network to handle"
- different from flow control (receiver controlled)!
- congestion indicators:
 - lost packets (buffer overflow at routers)
 - long delays (queueing in router buffers)

Flow vs congeston control

- (Flow control): the inability of the receiver accepting more data. Techniques addressing this problem are known as flow control, typically controlled by receiver.
- (Congestion control): the inability of the network delivering more data due to router buffer overflow. Techniques addressing this problem are known as congestion control, which include "sender-based flow control" (transmission rate) or dynamic routing (alternative routes).

Congestion collapse




History of TCP congestion control

- 1973: 3-way handshake introduced
- 1978: TCP/IP split into TCP and IP
- 1986: Internet suffered congestion collapses
- 1988: TCP (Tahoe) introduced congestion control
- 1990: fast retransmit and recovery added by TCP (Reno), supported by most hosts today.

TCP before Tahoe

- Upon connection established, a sender transmits full receiver window size segments, as fast as possible!
- Upon timeout, retransmit all unACK'ed segments immediately!
- The network is full of window-size segments.
- Without control, congestion collapse occurs!

TCP Congestion Control



sender limits transmission:

congestion window cwnd is dynamic, a kind of "sender-based flow control".

TCP sending rate:

 roughly: send cwnd bytes, wait RTT for ACKs, then send more bytes

rate
$$\approx \frac{\text{cwnd}}{\text{RTT}}$$
 bytes/sec

TCP sender self-clocking

- A sender's congestion window (cwnd) is adjusted automatically via returned ACKs.
- If the network is not congested, then every new ACK will trigger cwnd adjustment.
- The faster the new ACK arrives (i.e., short RTT), the more frequent cwnd is adjusted.

TCP: congestion detection

timeout:

- TimeOutInternval is calibrated upon every ACK received.
- When a timeout occurs, congestion is a likely cause.

triple duplicate ACK(n)s:

- Some data segments are lost but traffic is still getting through.
- Duplicate ACK(n)s are most likely caused by out-oforder segments; multiple segments after "n" have been sent but not ACK'ed.

TCP congestion control algorithm

- TCP (Reno) has 3 phases:
 - Slow Start (SS)
 - Congestion Avoidance (CA)
 - Fast Recovery (FR)
- **SS** is about "bandwidth probing", find out what is max. available capacity of the connection.
- CA is about "slowing down" when congestion is about to occur.
- Finally, FR is about "restoring optimal" transmission rate after receiving too many duplicate ACKs.

TCP: state variables and events

- In addition to cwnd (sender's congestion window), there is a ssthresh (threshold) triggering SS to CA phase change.
- cwnd and ssthresh are expressed in terms of MSS (max. segment size) bytes, typically 1460 bytes for MTU=1500 bytes.
- The following events trigger phase changes:
 - newACK: a new ACK just received
 - 3ACK(n): triple duplicate ACK(n)s
 - **timeout**: the **oldest** unACK'ed segment has just timeout.

TCP: initial conditions

- Upon connection setup, the receiver "advertises" its rwnd (max. free buffer space).
- Assume rwnd is always kept at its maximum (i.e., the receiver's host always picks up its data as fast as possible). Hence, we ignore flow control problems for now.
- To simplify our algorithm, we initialize cwnd to 2*64KB/MSS (will become clear later).
 (Note: TCP can send up to 64KB per IP datagram!)
- The algorithm starts in SS phase.

TCP Slow start

```
1 SS: // probing bandwidth
  ssthresh = cwnd/2;
   cwnd = 1;
                                                     one segment
                                          RT
   loop {
     event newACK:
                                                     two segments
 6
        cwnd += 1;
        transmit new segments;
 8
        if (cwnd==ssthresh) goto CA;
     event timeout:
                                                      four segments
10
        retransmit unACK'ed segments;
11
        goto SS;
12
     event 3ACK(n):
13
        retransmit segment @ n;
14
        goto FR;
15 }
                                                               time
```

Host A

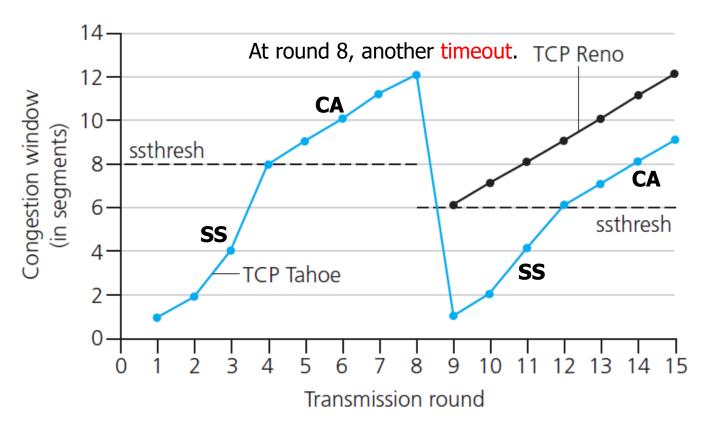
Host B

TCP Congestion avoidance

```
17 CA: // slowing down
18 loop {
19
     event newACK:
20
       cwnd += 1/cwnd;
21
       transmit new segments;
22 event timeout:
23
       goto SS;
24
     event 3ACK(n):
25
       retransmit segment @ n;
26
       goto FR;
27
```

TCP: SS to CA (example)

At round 0, a timeout occurs when cwnd=16; hence, ssthresh=16/2=8, and cwnd=1 at round 1.



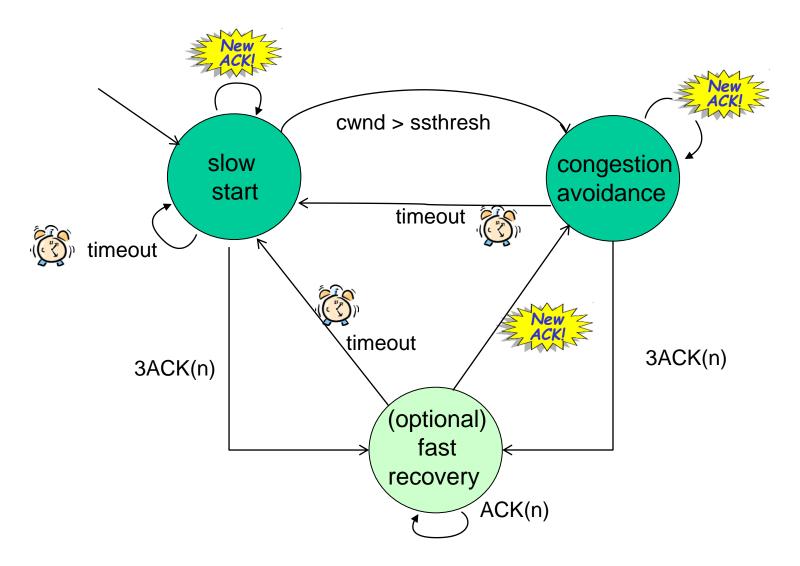
TCP (Reno) Fast recovery

```
29 FR: // restoring rate
30 ssthresh = cwnd/2;
31 \text{ cwnd} = \text{ssthresh} + 3;
32 loop{
33
      event newACK:
34
        goto CA;
35 | event timeout:
36
        goto SS;
      event ACK(n):
37
38
        cwnd += 1;
        retransmit segments @ n up
39
40
        to cwnd;
41
```

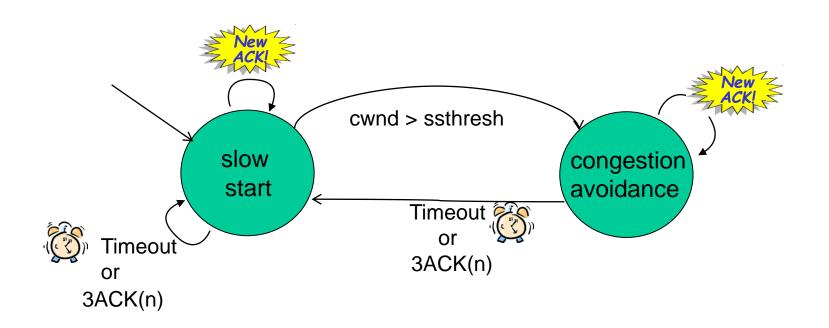
TCP (Tahoe)

- Fast recovery (FR optional) was introduced by TCP (Reno) around 1990, which is supported by most TCP installations today.
- However, TCP (Tahoe) doesn't support
 FR.
- Events timeout and 3ACK(n) are treated the same, as a segment loss; as a result, TCP (Tahoe) goes back to SS phase.

TCP (Reno): Congestion control



TCP (Tahoe): Congestion control



TCP: additive I multiplicative D

- transmission rate (cwnd) is increased until loss (timeout or 3ACKs) occurs
 - additive increase: increase cwnd by I

cwnd: TCP sender

 multiplicative decrease: set ssthresh=cwnd/2, and set cwnd=l

AIMD sawtooth behavior: probing for bandwidth

