

CSC361

Computer Networking

Mantis Cheng

Dept of Computer Science

Unit 2

Models of TCP/IP

Important Concepts

- IP Service Model
- UDP/TCP Service Model
- Multiplexing & Demultiplexing
- Client/Server Architecture
- Sockets API
- Python Sockets

What We Learned So Far

- IP is an **unreliable datagram delivery** service.
- TCP provides a **reliable bidirectional byte stream** service over IP.
- Most Internet applications are built on top of UDP/TCP using Berkeley **socket** API.
- **New** Internet application **functionalities** should only be added at the endpoints.
- **Abstraction** and **encapsulation** are key to the success of the 4-layer model of TCP/IP.

The IP Service Model (17:46)

Summary

- The Network (or IP) layer is the **most important** layer in the Internet.
- Each IP packet consists a **header** and a **payload**.
- Each IP **packet** is transmitted over the Link layer as a payload in a **frame**.
- The header of each IP packet contains both **source** and **destination** IP addresses.
- The Network layer only uses **IP addresses** to route each packet.

Summary (continued)

- Each router only provides a Network layer service, but **not** any (Transport/Application) layer above it.
- IP is **connectionless**; each packet is routed **independently**.
- IP is **unreliable**; it delivers datagrams with **best-effort** but **no** guarantee.
- To prevent packet delivery looping, each packet carries a **Time-To-Live** counter.

Summary (continued)

- IP may **fragment** long data segments into multiple packets.
- IP **checksums** its header to prevent delivery to the wrong address.
- IPv4 uses **32-bit** addresses. (We won't cover IPv6 in this course.)
- IP provides multiple **protocol ID** for demultiplexing to upper layer services (e.g., TCP, UDP, etc.)

UDP Service Model (6:41)

Summary

- User Datagram Protocol (UDP) is an **unreliable** datagram service built on top of IP (Network) layer.
- UDP is **connectionless**.
- UDP uses port numbers to **multiplex/demultiplex** datagrams to different applications.
- The UDP provides a **source** port number for the sender and a **destination** port number for the receiver.

Summary (continued)

- An IP address plus port number pair is known as a **socket**, a communication endpoint.
- When an application sends a UDP datagram, the UDP layer adds a **port number** for this application, and IP layer adds its **local host IP address** as source.
- The application **must** specify the receiver's socket (its port number and IP address).
- The port numbers identify the applications at both ends.

TCP Service Model (16:27)

Summary

- Transmission Control Protocol (TCP) is a **reliable, bidirectional, byte-stream** service built on top of IP (Network) layer.
- TCP is **connection-based**.
- TCP uses a **3-way handshake** to establish a connection.
- Both ends **must** agree upon an **initial sequence number** for each direction when a connection is made.

Summary (continued)

- A connection is two **continuous** streams of bytes in both directions, indexed by **sequence numbers**.
- A connection is closed with a **4-way tear-down**.
- Each TCP segment is checksummed.
- TCP uses **Flow** control to prevent receiver **buffer overflow**.
- All bytes are transmitted in **order reliably**.
- TCP uses **Congestion** control to prevent **Internet collapse**.

Summary (continued)

- Some port numbers are **well-known** (e.g., `TCP/80` is `HTTP`, `UDP/53` is `DNS`).
- The first 1024 ports are reserved.
- The remaining (64K-1K) ports are typically assigned **dynamically**.
- Each TCP connection is identified by two pairs of sockets ((`source IP address`, `source port #`), (`dest IP address`, `dest port #`)).

Wireshark Demo

(UDP, TCP, DNS)

(curl info.cern.ch)

Client-Server Model

(1:00:00)

**(watch the first 35 minutes of this
video at home)**

Summary

- A server is a process **waiting** for client's requests.
- A client is another process, on the same or different machine, making requests to the server.
- A server may serve **multiple** clients at the same time (*concurrently*); or it may serve **one** client's request at a time (*iteratively*).
- Iterative servers are suitable for **short service time** applications.
- Concurrent servers are suitable for **indeterminate service time** applications.

Summary (continued)

- Client/Server may use UDP or TCP for network communication.
- Typical internet applications use Berkeley Socket API.
- A `socket` = `TCP/UDP` + `IP address` + `Port number`.
- Two sockets identify two communication endpoints.
- UDP is an **unreliable datagram** service; while TCP is a **reliable bidirection byte-stream** service.

Summary (continued)

- DNS translates **hostnames** to **IP addresses**.
- A **domain** defines a collection of machines that share some common characteristics.
- A **Top-Level-Domain** (TLD) defines a top-level partition of all domain/host names in the Internet.
- DNS is a distributed database maintaining the mapping of hostnames to IP addresses.
- Watch [ICANN and the 7 Keys to the Internet \(5:03\)](#) to appreciate the security of DNS.

Socket Programming (34:00)

(watch the first 16 minutes)

Summary

- A `socket` = `TCP/UDP` + `IP address` + `Port number`.
- A port number corresponds to a **unique application id**.
- There are many **well-known** port numbers. Each port number is **16-bit** (64K ports) number.
- The [Berkeley socket API](#) includes: `socket`, `bind`, `connect`, `listen`, `accept`, `read`, `write`, and `close`.

Summary (continued)

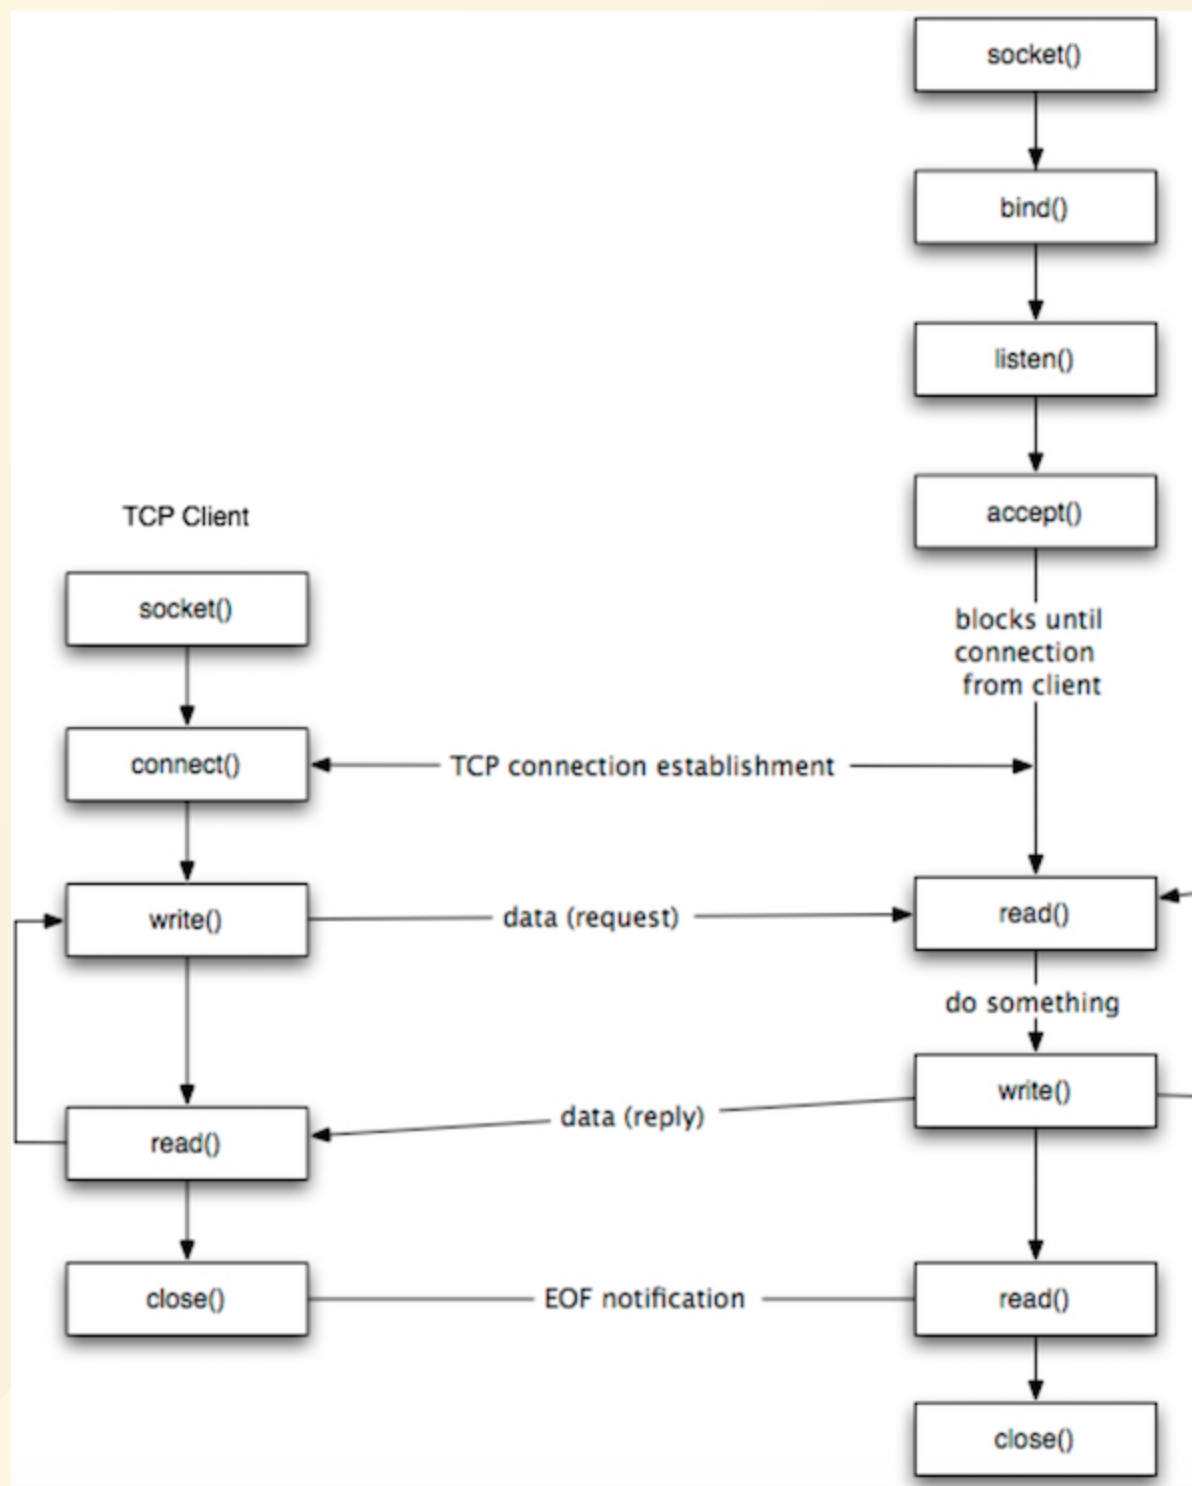
- **Server** side:

1. create a `socket` for accepting requests
2. `bind` the created socket to a port number
3. `listen` for any requests coming into the socket
4. `accept` one request from the socket
5. `receive/read` and process the request
6. `send/write` a response back to the client
7. `close` connection

Summary (continued)

- **Client** side:

1. create a `socket` for sending a request using UDP/TCP
2. `connect` its socket to the server's socket
3. `send/write` a request
4. `receive/read` a response



Summary (continued)

- `socket`: create a socket structure
- `connect`: establish a connection if needed
- `bind`: associate a socket to a port #
- `listen`: wait for requests
- `accept`: take one outstanding request
- `send/write`: transmit a request/response
- `receive/read`: receive a request/response
- `close`: close the connection if exist

Python Network **Programming Summary** **(PDF)**

**(study this summary and follow
the examples carefully)**

Python Network **Programming (21:38)**

(watch this video at home)

Summary

- It is very easy to use Python to create network applications using Berkeley sockets.
- The Python socket API is essentially **identical** to the standard C version.
- Our programming labs are all written in Python.
- Please follow the above examples to learn how to use UDP or TCP to create network applications.
- You will then understand the difference between UDP and TCP better afterwards.

The End